

```
!pip install transformers  
!pip install evaluate
```

[显示隐藏的输出项](#)

```
import torch  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
from datasets import load_dataset  
import transformers  
from transformers import AutoTokenizer, AutoModelForSequenceClassification, Trainer, TrainingArguments  
import evaluate  
from sklearn.metrics import confusion_matrix, classification_report
```

```
# =====  
# Evaluation Metrics  
# =====  
print(transformers.__version__)  
acc_metric = evaluate.load("accuracy")  
f1_metric = evaluate.load("f1")  
  
def compute_metrics(eval_pred):  
    logits, labels = eval_pred.predictions, eval_pred.label_ids  
    preds = np.argmax(logits, axis=1)  
    acc = acc_metric.compute(predictions=preds, references=labels)["accuracy"]  
    f1_macro = f1_metric.compute(predictions=preds, references=labels, average="macro")["f1"]  
    f1_weighted = f1_metric.compute(predictions=preds, references=labels, average="weighted")["f1"]  
    return {"accuracy": acc, "f1_macro": f1_macro, "f1_weighted": f1_weighted}
```

[显示隐藏的输出项](#)

```
# =====  
# Loading data  
# =====  
DATASET_NAME = "ag_news"  
raw = load_dataset(DATASET_NAME)  
split = raw["train"].train_test_split(test_size=0.08, seed=42)  
train_ds = split["train"]  
val_ds = split["test"]  
test_ds = raw["test"]  
  
label_names = train_ds.features["label"].names  
num_labels = len(label_names)
```

[显示隐藏的输出项](#)

```
# =====  
# Tokenization  
# =====  
def tokenize_datasets(model_ckpt, train_ds, val_ds, test_ds, max_length=128):  
    tokenizer = AutoTokenizer.from_pretrained(model_ckpt, use_fast=True)  
    def _tok(batch):  
        return tokenizer(batch["text"], truncation=True, padding='max_length', max_length=max_length)  
    t_train = train_ds.map(_tok, batched=True, batch_size=1000)  
    t_val = val_ds.map(_tok, batched=True, batch_size=1000)  
    t_test = test_ds.map(_tok, batched=True, batch_size=1000)  
    t_train.set_format(type="torch", columns=tokenizer.model_input_names + ["label"])  
    t_val.set_format(type="torch", columns=tokenizer.model_input_names + ["label"])  
    t_test.set_format(type="torch", columns=tokenizer.model_input_names + ["label"])  
    return tokenizer, t_train, t_val, t_test
```

```
# =====  
# Fine-tuning model  
# =====  
def fine_tune_resume(model_ckpt, outdir, tokenizer, train_dataset, val_dataset, test_dataset,  
                     num_labels, epochs, per_device_batch, lr, resume_from_checkpoint=None):  
    model = AutoModelForSequenceClassification.from_pretrained(model_ckpt, num_labels=num_labels)  
  
    training_args = TrainingArguments(  
        output_dir=outdir,  
        eval_strategy="epoch", # Automatic verification at each epoch  
        save_strategy="epoch", # Save checkpoints at each epoch  
        save_total_limit=5,  
        load_best_model_at_end=True,  
        metric_for_best_model="eval_accuracy",  
        per_device_train_batch_size=per_device_batch,
```

```

        per_device_eval_batch_size=per_device_batch,
        learning_rate=lr,
        weight_decay=0.01,
        num_train_epochs=epochs,
        fp16=True,
        logging_dir=f"{outdir}/logs",
        report_to="none"
    )

    trainer = Trainer(
        model=model,
        args=training_args,
        train_dataset=train_dataset,
        eval_dataset=val_dataset,
        compute_metrics=compute_metrics
    )

    # =====
    # determine whether to resume
    # =====
    if resume_from_checkpoint:
        print(f"Resuming training from checkpoint: {resume_from_checkpoint}")
    else:
        print("Starting training from scratch")

    trainer.train(resume_from_checkpoint=resume_from_checkpoint)

    # =====
    # Test set evaluation
    # =====
    pred = trainer.predict(test_dataset)
    print("Final Test metrics:", pred.metrics)

    return trainer, model, pred

```

```

# =====
# Confusion Matrix & Error Analysis
# =====
def plot_confusion_matrix(predictions, labels, label_names, title="Confusion Matrix"):
    cm = confusion_matrix(labels, predictions)
    plt.figure(figsize=(6, 5))
    sns.heatmap(cm, annot=True, fmt="d", xticklabels=label_names, yticklabels=label_names, cmap="Blues")
    plt.xlabel("Predicted")
    plt.ylabel("True")
    plt.title(title)
    plt.show()

def error_analysis(preds, labels, dataset, text_column="text", label_names=None, n_samples=10):
    import random

    errors = []

    # Text
    if text_column in dataset.column_names:
        texts = dataset[text_column]
    else:
        texts = [f"No {text_column}, sample #{i}" for i in range(len(dataset))]

    for i, (p, l) in enumerate(zip(preds, labels)):
        if p != l:
            errors.append({
                "text": texts[i],
                "true_label": label_names[1] if label_names else 1,
                "pred_label": label_names[p] if label_names else p
            })

    sample_errors = random.sample(errors, min(n_samples, len(errors)))

    for e in sample_errors:
        print(f"Text: {e['text']}")
        print(f"True: {e['true_label']} | Pred: {e['pred_label']}")
        print("-" * 50)

    return errors

```

```

# =====
# DistilBERT Fine-tuning
# =====
DISTIL_CKPT = "distilbert-base-uncased"

```

```

DISTIL_LR = 2e-5
DISTIL_BATCH = 32
DISTIL_EPOCHS = 3

tokenizer_distil, t_train_distil, t_val_distil, t_test_distil = tokenize_datasets(
    DISTIL_CKPT, train_ds, val_ds, test_ds
)

trainer_distil, model_distil, pred_distil = fine_tune_resume(
    DISTIL_CKPT, "./checkpoints/distilbert", tokenizer_distil,
    t_train_distil, t_val_distil, t_test_distil, num_labels,
    epochs=DISTIL_EPOCHS, per_device_batch=DISTIL_BATCH, lr=DISTIL_LR,
    resume_from_checkpoint=None # First Training
)

# Confusion Matrix & Error Analysis
distil_preds = np.argmax(pred_distil.predictions, axis=1)
plot_confusion_matrix(distil_preds, pred_distil.label_ids, label_names, title="DistilBERT Confusion Matrix")

```

tokenizer_config.json: 100%	48.0/48.0 [00:00<00:00, 2.01kB/s]
config.json: 100%	483/483 [00:00<00:00, 12.5kB/s]
vocab.txt: 100%	232k/232k [00:00<00:00, 3.08MB/s]
tokenizer.json: 100%	466k/466k [00:00<00:00, 7.39MB/s]
Map: 100%	110400/110400 [00:57<00:00, 1445.83 examples/s]
Map: 100%	9600/9600 [00:04<00:00, 3419.33 examples/s]
Map: 100%	7600/7600 [00:01<00:00, 5264.16 examples/s]
model.safetensors: 100%	268M/268M [00:05<00:00, 43.8MB/s]


Some weights of DistilBertForSequenceClassification were not initialized from the model checkpoint at distilbert-base-uncased and are newly You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Starting training from scratch

[10350/10350 18:10, Epoch 3/3]

Epoch	Training Loss	Validation Loss	Accuracy	F1 Macro	F1 Weighted
1	0.189900	0.182267	0.939271	0.938873	0.939144
2	0.135000	0.180944	0.942708	0.942408	0.942675
3	0.089800	0.202472	0.942500	0.942258	0.942517

Final Test metrics: {'test_loss': 0.1727038323879242, 'test_accuracy': 0.9423684210526316, 'test_f1_macro': 0.9423668324849056, 'test_f1_weighted': 0.9423668324849056}



errors_distil = error_analysis(distil_preds, pred_distil.label_ids, test_ds, label_names=label_names)

Text: Microsoft signs two Indian deals Microsoft is to form multi-million pound partnerships with two Indian software firms, and is expecte
True: Sci/Tech | Pred: Business

Text: Paris Tourists Search for Key to 'Da Vinci Code' (Reuters) Reuters - A funny thing happened on the way to the\ Mona Lisa. Visitors to
True: World | Pred: Sci/Tech

Text: Ali gives Iraq fighting chance The back of his shirt told the story last night at the Peristeri Olympic Boxing Hall.
True: Sports | Pred: World

Text: Is this the end of IT as we know it? Halsey Minor, CEO of hosted integration provider grand central Communications, has a powerful me
True: Business | Pred: Sci/Tech

Text: CA, Partners Move On As Kumar Faces Charges Concern over the fate of former Computer Associates International chairman and CEO Sanjay
True: Business | Pred: Sci/Tech

Text: Vodafone targets Japan with 3G offensive Vodafone has unveiled plans for 10 new third-generation handsets for Christmas to help shore
True: Business | Pred: Sci/Tech

Text: The battle for DR Congo's wildlife As donors pledge \\$40m to save DR Congo's wildlife, life in Virunga Park remains a battle ground.
True: Sci/Tech | Pred: World

Text: Labor situation deserves honest talk For a moment last week, President Bush escaped the White House spin chamber and was the plainspo
True: Business | Pred: World

Text: FDA OKs Scientist Publishing Vioxx Data (AP) AP - The Food and Drug Administration has given a whistle-blower scientist permission to
True: World | Pred: Sci/Tech

Text: Stocks Dip on Consumer Income Report News (AP) AP - An unsettling report on consumer incomes set off a spate of profit-taking on Wall
True: Business | Pred: World

```
# =====
# ModernBERT Fine-tuning
# =====
import torch._dynamo
torch._dynamo.config.suppress_errors = True

MODERN_CKPT = "answertotai/ModernBERT-base"
MODERN_LR = 3e-5
MODERN_BATCH = 16
MODERN_EPOCHS = 3

tokenizer_modern, t_train_modern, t_val_modern, t_test_modern = tokenize_datasets(
    MODERN_CKPT, train_ds, val_ds, test_ds
)

trainer_modern, model_modern, pred_modern = fine_tune_resume(
    MODERN_CKPT, "./checkpoints/modernbert", tokenizer_modern,
    t_train_modern, t_val_modern, t_test_modern, num_labels,
    epochs=MODERN_EPOCHS, per_device_batch=MODERN_BATCH, lr=MODERN_LR,
    resume_from_checkpoint=None # First Training
)

# Confusion Matrix & Error Analysis
modern_preds = np.argmax(pred_modern.predictions, axis=1)
plot_confusion_matrix(modern_preds, pred_modern.label_ids, label_names, title="ModernBERT Confusion Matrix")
```

```

tokenizer_config.json:      20.8k/? [00:00<00:00, 1.33MB/s]
tokenizer.json:      2.13M/? [00:00<00:00, 29.6MB/s]
special_tokens_map.json: 100%                                     694/694 [00:00<00:00, 70.1kB/s]

Map: 100%                                         110400/110400 [00:27<00:00, 3987.21 examples/s]
Map: 100%                                         9600/9600 [00:02<00:00, 3482.09 examples/s]
Map: 100%                                         7600/7600 [00:02<00:00, 4238.01 examples/s]

config.json:     1.19k/? [00:00<00:00, 98.1kB/s]


model.safetensors: 100%                                     599M/599M [01:08<00:00, 14.4MB/s]

Some weights of ModernBertForSequenceClassification were not initialized from the model checkpoint at answerdotai/ModernBERT-base and are not trained by this script.
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
Starting training from scratch
W0913 21:55:14.971000 413 torch/_inductor/utils.py:1436] [1/0_1] Not enough SMs to use max_autotune_gemm mode
[20700/20700 1:05:30, Epoch 3/3]

```

Epoch	Training Loss	Validation Loss	Accuracy	F1 Macro	F1 Weighted
1	0.174800	0.179104	0.942500	0.942170	0.942427
2	0.089600	0.203006	0.947604	0.947357	0.947595
3	0.038400	0.350692	0.945625	0.945384	0.945650

Final Test metrics: {'test_loss': 0.20180293917655945, 'test_accuracy': 0.9460526315789474, 'test_f1_macro': 0.9459716717408502, 'test_f1_weighted': 0.9459716717408502}



```
errors_modern = error_analysis(modern_preds, pred_modern.label_ids, test_ds, label_names=label_names)
```

- Text: TheStreet.com May Be Up for Sale -- Report (Reuters) Reuters - TheStreet.com Inc. , the\financial news and commentary Web site, may be up for sale.

True: Sci/Tech | Pred: Business
- Text: China Reports Births of Two Giant Pandas (AP) AP - For pandas, it's practically a baby boom. Two giant pandas were born this week, an AP report said.

True: Sci/Tech | Pred: World
- Text: Court: File-swapping software not liable for copyright violations The makers of two leading file-sharing programs are not legally liable for copyright violations.

True: Sci/Tech | Pred: Business
- Text: Amazon #39;s UK DVD Rentals Might Presage War with Netflix Analysts said Amazon #39;s launch of the DVD rental service in the UK might presage a war between the two companies.

True: Business | Pred: Sci/Tech
- Text: Adobe updates RAW plug-in with digital negative format Adobe has updated Photoshop #39;s support for digital cameras #39; RAW image format.

True: Business | Pred: Sci/Tech
- Text: PeopleSofts big bash See you next year in Las Vegas , proclaimed a marquee at the PeopleSoft user conference in San Francisco in late October.

True: Business | Pred: Sci/Tech
- Text: Ask Jeeves retools search engine in bid to catch Google, Yahoo SAN FRANCISCO -- Hoping to emerge from the shadow of its more popular search engines, Ask Jeeves is retooling its search engine.

True: Business | Pred: Sci/Tech
- Text: Retailers Looking to Move Plasma TV's (AP) AP - Hanging stockings by the chimney with care? Retailers hope that St. Nicholas soon will bring them more customers.

True: Business | Pred: Sci/Tech
- Text: English 'world language' forecast A third of people on the planet will be learning English in the next decade, \says a report.

True: World | Pred: Sci/Tech
- Text: Hall of Shame Hall of Fame We spotlight people and products that pester us...and the heroes saving us from annoyances.

```

# %% Evaluation and comparison (accuracy / macro F1 / per-class)
# Distil
from sklearn.metrics import precision_recall_fscore_support
val_pred_distil = trainer_distil.predict(t_val_distil)
y_val_true = val_pred_distil.label_ids
y_val_pred = np.argmax(val_pred_distil.predictions, axis=1)
y_true = pred_distil.label_ids
y_pred_distil = np.argmax(pred_distil.predictions, axis=1)

# Modern
val_pred_modern = trainer_modern.predict(t_val_modern)
y_val_true_mod = val_pred_modern.label_ids
y_val_pred_mod = np.argmax(val_pred_modern.predictions, axis=1)
y_true_mod = pred_modern.label_ids      # Should be consistent with y_true (same test set)
y_pred_modern = np.argmax(pred_modern.predictions, axis=1)

from sklearn.metrics import accuracy_score, f1_score
summary = pd.DataFrame([
    {
        "model": "DistilBERT",
        "val_accuracy": accuracy_score(y_val_true, y_val_pred),
        "val_f1_macro": f1_score(y_val_true, y_val_pred, average="macro"),
        "test_accuracy": accuracy_score(y_true, y_pred_distil),
        "test_f1_macro": f1_score(y_true, y_pred_distil, average="macro"),
    },
    {
        "model": "ModernBERT",
        "val_accuracy": accuracy_score(y_val_true_mod, y_val_pred_mod),
        "val_f1_macro": f1_score(y_val_true_mod, y_val_pred_mod, average="macro"),
        "test_accuracy": accuracy_score(y_true_mod, y_pred_modern),
        "test_f1_macro": f1_score(y_true_mod, y_pred_modern, average="macro"),
    }
])
display(summary)

# per-class metrics
def per_class_table(y_true, y_pred, labels):
    p, r, f, s = precision_recall_fscore_support(y_true, y_pred, labels=list(range(len(labels))))
    df = pd.DataFrame({"label": labels, "precision": p, "recall": r, "f1": f, "support": s})
    return df

df_distil_cls = per_class_table(y_true, y_pred_distil, label_names)
df_modern_cls = per_class_table(y_true_mod, y_pred_modern, label_names)
display(df_distil_cls)
display(df_modern_cls)

```

	model	val_accuracy	val_f1_macro	test_accuracy	test_f1_macro	
0	DistilBERT	0.942708	0.942408	0.942368	0.942367	
1	ModernBERT	0.947604	0.947357	0.946053	0.945972	

	label	precision	recall	f1	support	
0	World	0.952830	0.956842	0.954832	1900	
1	Sports	0.989424	0.984737	0.987075	1900	
2	Business	0.925081	0.896842	0.910743	1900	
3	Sci/Tech	0.903012	0.931053	0.916818	1900	

	label	precision	recall	f1	support	
0	World	0.955091	0.962632	0.958847	1900	
1	Sports	0.986352	0.988947	0.987648	1900	
2	Business	0.926106	0.903684	0.914758	1900	
3	Sci/Tech	0.916407	0.928947	0.922635	1900	

后续
步
骤：

[使用 summary 生成代码](#)

[查看推荐的图表](#)

[New interactive sheet](#)

[使用 df_distil_cls 生成代码](#)

[查看推荐的图表](#)

[New interactive sheet](#)

```

# %% Confusion Matrix (normalized)
def plot_cm(y_true, y_pred, labels, title):
    cm = confusion_matrix(y_true, y_pred)
    cmn = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]



```

```

plt.figure(figsize=(7, 5))
sns.heatmap(cm, annot=True, fmt=".2f", xticklabels=labels, yticklabels=labels)
plt.xlabel("Pred")
plt.ylabel("True")
plt.title(title)
plt.show()

plot_cm(y_true, y_pred_distil, label_names, "DistilBERT - normalized CM")
plot_cm(y_true_mod, y_pred_modern, label_names, "ModernBERT - normalized CM")

```



```

import torch.nn.functional as F
# %% Error analysis: save misclassified samples and increase confidence level of misclassified samples
def make_error_df(pred_output, tokenizer, raw_test_ds):
    logits = pred_output.predictions
    probs = F.softmax(torch.tensor(logits), dim=1).numpy()
    preds = np.argmax(logits, axis=1)
    labels = pred_output.label_ids
    texts = raw_test_ds["text"] # Original text
    df = pd.DataFrame({"text": texts, "true_id": labels, "pred_id": preds})
    df["true"] = df["true_id"].map(lambda x: label_names[x])
    df["pred"] = df["pred_id"].map(lambda x: label_names[x])
    # Confidence
    df["pred_conf"] = np.max(probs, axis=1)
    df_err = df[df["true_id"] != df["pred_id"]].copy()
    df_err = df_err.sort_values("pred_conf", ascending=False)
    return df_err

err_distil = make_error_df(pred_distil, tokenizer_distil, test_ds)
err_modern = make_error_df(pred_modern, tokenizer_modern, test_ds)

# Save the top 200 error samples as csv
err_distil.head(200).to_csv("errors_distil_top200.csv", index=False)

```

```

err_distil.head(200).to_csv('errors_distil_top200.csv', index=False)
print("Saved error CSVs.")

# Print several high-confidence but wrong examples
print("DistilBERT High-Confidence Misclassification Examples:")
display(err_distil.head(5)[["text", "true", "pred", "pred_conf"]])
print("ModernBERT High-Confidence Misclassification Examples:")
display(err_modern.head(5)[["text", "true", "pred", "pred_conf"]])

# %% Counting the types and number of errors made by DistilBERT
print("DistilBERT Error type statistics:")
display(err_distil.groupby(['true', 'pred']).size().sort_values(ascending=False).to_frame(name='count'))

# %% Counting the types and number of errors made by ModernBERT
print("ModernBERT Error type statistics:")
display(err_modern.groupby(['true', 'pred']).size().sort_values(ascending=False).to_frame(name='count'))

# %% Compare the misclassified sets of two models (intersection/difference)
set_distil_err_texts = set(err_distil["text"])
set_modern_err_texts = set(err_modern["text"])
unique_distil = set_distil_err_texts - set_modern_err_texts
unique_modern = set_modern_err_texts - set_distil_err_texts
common_err = set_distil_err_texts & set_modern_err_texts
print("DistilBERT Unique Misclassifications:", len(unique_distil))
print("ModernBERT Unique Misclassifications:", len(unique_modern))
print("Common Misclassifications:", len(common_err))

# Printing Example
print("DistilBERT Unique Examples (2 samples):")
for i, t in enumerate(list(unique_distil)[:2]): print(i+1, t)
print("ModernBERT Unique Examples (2 samples):")
for i, t in enumerate(list(unique_modern)[:2]): print(i+1, t)

```

Saved error CSVs.

DistilBERT High-Confidence Misclassification Examples:

		text	true	pred	pred_conf	grid icon
						bar chart icon
2126	Bush Scraps Most U.S. Sanctions on Libya (Reut...	Sci/Tech	World	0.999053		
912	Bryant Makes First Appearance at Trial (AP) AP...	Sci/Tech	Sports	0.999017		
1269	Philippines mourns dead in Russian school sieg...	Sci/Tech	World	0.998799		
7228	Indonesian diplomats asked to help improve RI ...	Business	World	0.998725		
1197	World briefs LONDON - A man wielding a machete...	Business	World	0.998257		

ModernBERT High-Confidence Misclassification Examples:

		text	true	pred	pred_conf	grid icon
						bar chart icon
7467	Daily Mail hits back at Blunkett The Daily Mai...	World	Sports	0.999990		
2660	Gas prices up 5 cents after Hurricane Ivan CAM...	Sci/Tech	Business	0.999981		
1197	World briefs LONDON - A man wielding a machete...	Business	World	0.999975		
4540	Miss Peru takes Miss World crown Twenty-year-o...	Sports	World	0.999951		
3077	Brewers buyer expected to step out of the shad...	Business	Sports	0.999946		

DistilBERT Error type statistics:

	count	grid icon
true	pred	bar chart icon
Business	Sci/Tech	148
Sci/Tech	Business	93
Business	World	43

```
# %% Install word cloud library
!pip install wordcloud
```

```
RSci/Tech atWorld satisfied: wordcloud in /usr/local/lib/python3.12/dist-packages (1.9.4)
Requirement already satisfied: numpy>=1.6.1 in /usr/local/lib/python3.12/dist-packages (from wordcloud) (2.0.2)
RSci/Tech atWorld satisfied: pillow in /usr/local/lib/python3.12/dist-packages (from wordcloud) (11.3.0)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (from wordcloud) (3.10.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib>wordcloud) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib>wordcloud) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib>wordcloud) (4.59.2)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib>wordcloud) (1.4.9)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib>wordcloud) (25.0)
RBusiness atSports satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib>wordcloud) (3.2.3)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.12/dist-packages (from matplotlib>wordcloud) (2.9.0.post0)
RSci/Tech atSports satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.7>matplotlib>wordcloud) (1.17.0)
...

```

```
# %% Generate a word cloud for a specific error type
```

```
import matplotlib.pyplot as plt
from wordcloud import WordCloud

Sci/Tech Business 95

def plot_wordcloud(err_df, true_label, pred_label, model_name):

    target_errors_df = err_df[(err_df['true'] == true_label) & (err_df['pred'] == pred_label)]

    if not target_errors_df.empty:
        text_corpus = ' '.join(target_errors_df['text'].tolist())
        wordcloud = WordCloud(width=800, height=400, background_color='white').generate(text_corpus)

        plt.figure(figsize=(10, 5))
        plt.imshow(wordcloud, interpolation='bilinear')
        plt.axis('off')
        plt.title(f'Word Cloud for {true_label} → {pred_label} Misclassifications ({model_name})')
        plt.show()

        print(f"Found {len(target_errors_df)} examples where '{true_label}' was misclassified as '{pred_label}'.")
    else:
        print(f"No error samples found for {true_label} → {pred_label} in {model_name}.")
```


ModernBERT Unique Misclassifications: 106

```
# DistilBERT: Business → Sci/Tech
plot_wordcloud(err_distil, "Business", "Sci/Tech", "DistilBERT")
```

2 E.U. Regulators Approve Oracle's PeopleSoft Bid European Union antitrust regulators on Tuesday cleared Oracle Corp.'s hostile \$7.7 billion ModernBERT Unique Examples (2 samples):

1 Mazu scores VC lifeblood from Symantec, others Mazu Networks, the Cambridge, Massachusetts, network intrusion prevention system (IPS) technology 2 SBC calling on cable SBC is teaming up with Microsoft to provide consumers with a new way to view television -- a move that puts it in di

Word Cloud for Business -> Sci/Tech Misclassifications (DistilBERT)



Found 148 examples where 'Business' was misclassified as 'Sci/Tech'.

```
# DistilBERT: Sci/Tech -> Business  
plot_wordcloud(err_distil, "Sci/Tech", "Business", "DistilBERT")
```

Word Cloud for Sci/Tech -> Business Misclassifications (DistilBERT)



Found 93 examples where 'Sci/Tech' was misclassified as 'Business'.

```
# ModernBERT: Business -> Sci/Tech  
plot_wordcloud(err_modern, "Business", "Sci/Tech", "ModernBERT")
```

Word Cloud for Business -> Sci/Tech Misclassifications (ModernBERT)



Found 130 examples where 'Business' was misclassified as 'Sci/Tech'.
Found 0 texts containing 'Swill'.

```
# ModernBERT: Sci/Tech -> Business  
plot_wordcloud(err_modern, "Sci/Tech", "Business", "ModernBERT")
```

Word Cloud for Sci/Tech -> Business Misclassifications (ModernBERT)