

assignment_3

October 24, 2025

```
[ ]: # === IMPORTS AND CONFIG ===
from transformers import AutoTokenizer, AutoModelForCausalLM, □
    ↪AutoModelForSeq2SeqLM
from datasets import load_dataset
import evaluate
import torch
from tqdm.auto import tqdm
import pandas as pd
import os, json

[ ]: # === CONFIGURATIONS ===
DEVICE = 'cuda' if __import__('torch').cuda.is_available() else 'cpu'
SMOL_MODEL = 'HuggingFaceTB/SmolLM-135M-Instruct'
PEGASUS_MODEL = 'google/pegasus-cnn_dailymail'

MAX_EXAMPLES = 1000 # Use a subset for faster evaluation
MAX_TOKENS = 90 # Max new tokens for summary
MIN_LENGTH = 30 # Min length for summary

[ ]: # === DATA LOADING ===
def load_cnn_dailymail(split='test', max_examples=None):
    ds = load_dataset('cnn_dailymail', '3.0.0', split=split)
    if max_examples is not None:
        ds = ds.select(range(max_examples))
    return ds

[ ]: # === MODEL LOADERS ===
def load_smol(model_name=SMOL_MODEL):
    tok = AutoTokenizer.from_pretrained(model_name, use_fast=True, □
        ↪trust_remote_code=True, chat_template=None)
    tok.padding_side = 'left' # Left padding to avoid decoder-only warnings
    model = AutoModelForCausalLM.
        ↪from_pretrained(model_name, trust_remote_code=True).to(DEVICE)
    # ensure pad token
    if tok.pad_token is None:
        tok.add_special_tokens({'pad_token': '[PAD]'})
    model.resize_token_embeddings(len(tok))
```

```

    return tok, model

def load_pegasus(model_name=PEGASUS_MODEL):
    tok = AutoTokenizer.from_pretrained(model_name)
    model = AutoModelForSeq2SeqLM.from_pretrained(model_name).to(DEVICE)
    return tok, model

```

```
[ ]: # === PROMPT DEFINITIONS ===

def truncate_article(text, tokenizer, max_tokens=512):
    """Truncates text to max_tokens to fit in context."""
    tokens = tokenizer.encode(text, truncation=True, max_length=max_tokens)
    return tokenizer.decode(tokens, skip_special_tokens=True)

def prompt_instruction(article, tokenizer=None):
    """Creates a zero-shot instruction prompt."""
    if tokenizer:
        article = truncate_article(article, tokenizer)
    return (
        "You are an advanced news summarization assistant.\n"
        "Write a detailed, factual summary of the following article in 3-4"
        "complete sentences.\n"
        "Focus on the key events, people, and outcomes. Avoid repetition or"
        "speculation.\n\n"
        f"Article:\n{article.strip()}\n\n"
        "Summary:"
    )

# Define few-shot examples (using one representative example)
FEW_SHOT_EXAMPLES = [
(
    "A powerful winter storm has blanketed much of the northeastern United"
    "States in heavy snow, disrupting travel and closing schools across several"
    "states. "
    "New York City received more than eight inches of snow overnight, while"
    "parts of Massachusetts and Connecticut reported over a foot. "
    "Thousands of flights were canceled or delayed, leaving travelers"
    "stranded at major airports. "
    "Commuters faced treacherous road conditions, and authorities urged"
    "residents to stay home unless absolutely necessary. "
    "In Boston, city officials declared a snow emergency and deployed"
    "hundreds of plows to clear main roads. "
    "Utility companies reported widespread power outages as strong winds"
    "knocked down trees and power lines. "
)
```

```

        "Amtrak suspended most regional train services until conditions
↳improved. "
        "Meteorologists say the storm is expected to weaken by Thursday
↳afternoon but warned of another system developing early next week. "
        "Despite the disruption, some residents took advantage of the snowfall
↳to go sledding and build snowmen in city parks. "
        "Officials reminded the public to check on elderly neighbors and keep
↳emergency supplies at home in case of extended outages.",

        "Summary: A major winter storm brought heavy snow and high winds to the
↳U.S. Northeast, disrupting travel and shutting down schools. "
        "More than a foot of snow fell in some areas, causing widespread flight
↳cancellations and power outages. "
        "Officials declared emergencies, urged residents to stay indoors, and
↳warned of another storm approaching next week. "
        "The system is expected to weaken by Thursday afternoon as cleanup
↳efforts continue."
    )
]

```

```

def prompt_few_shot(article, tokenizer=None):
    """Creates a one-shot prompt."""
    if tokenizer:
        article = truncate_article(article, tokenizer)

    example_article, example_summary = FEW_SHOT_EXAMPLES[0]
    return (
        "You are an advanced summarization assistant.\n"
        "Here are examples of how to summarize news articles in 3-4 sentences:
↳\n\n"
        f"Example article:{example_article}\n{example_summary}\n\n"
        f"Now summarize the following article in 3-4 sentences:
↳\n{article}\n\nSummary:"
    )

```

```

[ ]: # === GENERATION WRAPPERS ===
def strip_prefix(prompt, text):
    """Removes the prompt from the start of the generated text."""
    p = prompt.strip()
    t = text.strip()
    if not p:
        return t
    if t.startswith(p):
        return t[len(p):].strip()
    return t

```

```

def generate_smol(tokenizer, model, prompts, strategy, max_new_tokens=90, ↴
    ↵min_length=50, batch_size=8):
    """Handles generation for the decoder-only SmoLLM."""
    model.eval()
    outputs = []

    for i in range(0, len(prompts), batch_size):
        batch = prompts[i:i+batch_size]
        inputs = tokenizer(batch, return_tensors='pt', padding=True, ↴
            ↵truncation=True, max_length=1024).to(DEVICE)

        with torch.no_grad():
            gen = model.generate(
                **inputs,
                max_new_tokens=max_new_tokens,
                min_length=min_length,
                **strategy
            )

        texts = tokenizer.batch_decode(gen, skip_special_tokens=True)

    # Handle cases where num_return_sequences > 1
    num_return = strategy.get('num_return_sequences', 1)
    if num_return > 1:
        grouped = [texts[j*num_return:(j+1)*num_return] for j in ↴
            range(len(batch))]
        texts = [g[0] for g in grouped]
    return outputs

def generate_pegasus(tokenizer, model, articles, max_new_tokens=90, ↴
    ↵min_length=50, strategy=None, batch_size=8):
    """Handles generation for the Seq2Seq PEGASUS model."""
    model.eval()
    strategy = strategy or {}
    outputs = []

    for i in range(0, len(articles), batch_size):
        batch = articles[i:i+batch_size]
        inputs = tokenizer(batch, return_tensors='pt', padding=True, ↴
            ↵truncation=True, max_length=1024).to(DEVICE)

        with torch.no_grad():
            gen = model.generate(
                **inputs,
                max_new_tokens=max_new_tokens,
                min_length=min_length,

```

```

        **strategy
    )

texts = tokenizer.batch_decode(gen, skip_special_tokens=True)

# Handle num_return_sequences > 1
num_return = strategy.get('num_return_sequences', 1)
if num_return > 1:
    texts = [texts[j * num_return] for j in range(len(batch))]

outputs.extend([t.strip() for t in texts])
return outputs

```

```

[ ]: # === EVALUATION HELPER ===
rouge = evaluate.load('rouge')

def compute_rouge(preds, refs):
    """Computes ROUGE scores."""
    result = rouge.compute(predictions=preds, references=refs)
    # Helper to extract f-measure
    def _get(k):
        v = result.get(k)
        if hasattr(v, 'mid'):
            return v.mid.fmeasure
        return v
    return {'rouge1': _get('rouge1'), 'rouge2': _get('rouge2'), 'rougeL':_get('rougeL')}

```

```

[ ]: # === EXPERIMENT RUNNER ===
def run_experiments(
    max_examples=MAX_EXAMPLES,
    min_length=MIN_LENGTH,
    save_outputs_path='results.jsonl',
    mode='smol'
):
    """Runs all experiments for a given mode ('smol' or 'pegasus')."""
    # Load data
    ds = load_cnn_dailymail('test', max_examples=max_examples)
    articles = [ex['article'] for ex in ds]
    refs = [ex['highlights'] for ex in ds]

    experiments = []

    # --- Task 2: SmollM Experiments ---
    if mode == 'smol':
        print('Loading SmollM...')
        tok_smol, model_smol = load_smol()

```

```

# Define the 4 strategies
all_experiments = [
    ("E1: Instruction+Greedy", lambda: [prompt_instruction(a, tok_smol) for a in articles], {'num_beams': 1, 'do_sample': False, 'no_repeat_ngram_size': 3}),
    ("E2: Instruction+Beam", lambda: [prompt_instruction(a, tok_smol) for a in articles], {'num_beams': 2, 'do_sample': False, 'early_stopping': True, 'no_repeat_ngram_size': 3}),
    ("E3: FewShot+Beam", lambda: [prompt_few_shot(a, tok_smol) for a in articles], {'num_beams': 2, 'do_sample': False, 'early_stopping': True, 'no_repeat_ngram_size': 3}),
    ("E4: FewShot+TopP", lambda: [prompt_few_shot(a, tok_smol) for a in articles], {'do_sample': True, 'num_beams': 1, 'top_p': 0.95, 'temperature': 0.7, 'no_repeat_ngram_size': 3}),
]

# Run each SmolM experiment
for name, prompt_fn, strategy in all_experiments:
    print(f'Running {name}...')
    prompts = prompt_fn()
    preds = generate_smol(tok_smol, model_smol, prompts, strategy, max_new_tokens=MAX_TOKENS, min_length=MIN_LENGTH)
    scores = compute_rouge(preds, refs)
    experiments.append({'name': name, 'scores': scores, 'preds': preds})
    torch.cuda.empty_cache()

# --- Task 3: PEGASUS Comparison ---
elif mode == 'pegasus':
    print('Loading PEGASUS...')
    tok_peg, model_peg = load_pegasus()
    # Define PEGASUS strategy
    strategy_peg = {'num_beams': 8, 'length_penalty': 0.8}
    preds_peg = generate_pegasus(tok_peg, model_peg, articles, max_new_tokens=MAX_TOKENS, min_length=MIN_LENGTH, strategy=strategy_peg)
    scores_peg = compute_rouge(preds_peg, refs)
    experiments.append({'name': 'PEGASUS', 'scores': scores_peg, 'preds': preds_peg})
    torch.cuda.empty_cache()

else:
    raise ValueError("mode must be 'smol' or 'pegasus'")

# --- RESULTS & SAVING ---
# Collate results into a DataFrame
rows = [{method: exp['name'], **exp['scores']} for exp in experiments]

```

```

df = pd.DataFrame(rows).set_index('method')

# Save raw outputs
with open(save_outputs_path, 'w', encoding='utf-8') as fh:
    for i in range(len(articles)):
        rec = {'id': ds[i]['id'], 'article': articles[i], 'reference': refs[i]}
        for exp in experiments:
            rec[exp['name']] = exp['preds'][i]
        fh.write(json.dumps(rec, ensure_ascii=False) + '\n')

return experiments, df

```

```

[ ]: # === MAIN EXECUTION ===
if __name__ == '__main__':
    # Step 1: Run SmoLLM experiments
    torch.cuda.empty_cache()
    res_smol, df_smol = run_experiments(mode='smol', max_examples=MAX_EXAMPLES,
                                          min_length=MIN_LENGTH, save_outputs_path='smol_results.jsonl')
    print(df_smol)
    df_smol.to_csv('smol_rouge_scores.csv')

    # Step 2: Run PEGASUS experiment
    torch.cuda.empty_cache()
    res_peg, df_peg = run_experiments(mode='pegasus',
                                        max_examples=MAX_EXAMPLES, min_length=MIN_LENGTH,
                                        save_outputs_path='pegasus_results.jsonl')
    print(df_peg)
    df_peg.to_csv('pegasus_rouge_scores.csv')

    # Step 3: Combine and save all results
    #df_smol = pd.read_csv('smol_rouge_scores.csv')
    df_total = pd.concat([df_smol, df_peg], ignore_index=False)
    print(df_total)
    df_total.to_csv('summary_rouge_scores.csv')

```

Loading SmoLLM...

Running E1: Instruction+Greedy...

Running E2: Instruction+Beam...

Running E3: FewShot+Beam...

Running E4: FewShot+TopP...

	rouge1	rouge2	rougeL
method			
E1: Instruction+Greedy	0.136200	0.016615	0.103861
E2: Instruction+Beam	0.137873	0.015332	0.103841
E3: FewShot+Beam	0.147683	0.017549	0.107364
E4: FewShot+TopP	0.148044	0.015964	0.106671

Loading PEGASUS...

```
model.safetensors: 0% | 0.00/2.28G [00:00<?, ?B/s]

Some weights of PegasusForConditionalGeneration were not initialized from the
model checkpoint at google/pegasus-cnn_dailymail and are newly initialized:
['model.decoder.embed_positions.weight', 'model.encoder.embed_positions.weight']
You should probably TRAIN this model on a down-stream task to be able to use it
for predictions and inference.

      rouge1  rouge2    rougeL
method
PEGASUS  0.321738  0.1335  0.237964
                  rouge1    rouge2    rougeL
method
E1: Instruction+Greedy  0.136200  0.016615  0.103861
E2: Instruction+Beam    0.137873  0.015332  0.103841
E3: FewShot+Beam        0.147683  0.017549  0.107364
E4: FewShot+TopP         0.148044  0.015964  0.106671
PEGASUS                 0.321738  0.133500  0.237964
```