

```
# %%
# !pip install transformers[torch] datasets evaluate accelerate optuna scikit-learn pandas matplotlib -q
!pip install evaluate optuna
#!pip install -U transformers bitsandbytes
```

[显示隐藏的输出项](#)

```
# %%
from datasets import load_dataset
from transformers import AutoTokenizer, AutoModelForSequenceClassification, TrainingArguments, Trainer, pipeline
import torch, numpy as np, evaluate, torch.nn as nn, torch.nn.functional as F, pandas as pd, matplotlib.pyplot as plt
from torch.quantization import quantize_dynamic
from time import perf_counter
from pathlib import Path
import optuna

# %%
clinc = load_dataset("clinc_oos", "plus")
teacher_ckpt = "answerdotai/ModernBERT-large"
student_ckpt = "bert-base-uncased"
intents = clinc["test"].features["intent"]
id2label = {i: label for i, label in enumerate(intents.names)}
label2id = {label: i for i, label in enumerate(intents.names)}
num_labels = intents.num_classes

teacher_tokenizer = AutoTokenizer.from_pretrained(teacher_ckpt)
def tokenize_text(batch): return teacher_tokenizer(batch["text"], truncation=True)
clinc_enc = clinc.map(tokenize_text, batched=True, remove_columns=["text"])
clinc_enc = clinc_enc.rename_column("intent", "labels")

student_tokenizer = AutoTokenizer.from_pretrained(student_ckpt)
def tokenize_text_student(batch): return student_tokenizer(batch["text"], truncation=True)
clinc_enc_student = clinc.map(tokenize_text_student, batched=True, remove_columns=["text"])
clinc_enc_student = clinc_enc_student.rename_column("intent", "labels")

accuracy_score = evaluate.load("accuracy")
f1_metric = evaluate.load("f1")
def compute_metrics(pred):
    # 'pred' is an EvalPrediction object
    predictions, labels = pred
    predictions = np.argmax(predictions, axis=1) # Convert logits to class IDs

    # Compute Accuracy
    accuracy_result = accuracy_score.compute(predictions=predictions, references=labels)

    # Compute Macro F1 Score
    # For multi-class classification, you must specify the 'average' parameter.
    f1_result = f1_metric.compute(
        predictions=predictions,
        references=labels,
        average="macro" # <-- This calculates the MACRO F1 score
    )

    # Combine results into a single dictionary
    # The keys must match what you are expecting, e.g., 'eval_macro_f1'
    return {
        "accuracy": accuracy_result["accuracy"],
        "macro_f1": f1_result["f1"], # Return 'f1' key from the compute result
    }
```

```

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it as secret in
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
    warnings.warn(
README.md:      24.0k/? [00:00<00:00, 719kB/s]

plus/train-00000-of-00001.parquet: 100%                                312k/312k [00:00<00:00, 347kB/s]
plus/validation-00000-of-00001.parquet: 100%                            77.8k/77.8k [00:00<00:00, 152kB/s]
plus/test-00000-of-00001.parquet: 100%                               136k/136k [00:00<00:00, 444kB/s]
Generating train split: 100%                                         15250/15250 [00:00<00:00, 124188.45 examples/s]
Generating validation split: 100%                           3100/3100 [00:00<00:00, 47332.88 examples/s]
Generating test split: 100%                                         5500/5500 [00:00<00:00, 90886.31 examples/s]
tokenizer_config.json:      20.8k/? [00:00<00:00, 518kB/s]
tokenizer.json:        2.13M/? [00:00<00:00, 30.6MB/s]
special_tokens_map.json: 100%                                         694/694 [00:00<00:00, 15.8kB/s]

```

```
Map: 100%                                         15250/15250 [00:00<00:00, 5947.35 examples/s]
```

```

teacher_training_args = TrainingArguments(
    output_dir="/kaggle/working/ModernBERT-large-finetuned-clinc",
    eval_strategy="epoch",
    num_train_epochs=5,
    learning_rate=2e-5,
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
    weight_decay=0.01,
    logging_steps=100,
    save_strategy="epoch",
    load_best_model_at_end=True,
    metric_for_best_model="accuracy",      # 确定最佳模型的评价指标
    greater_is_better=True,
    report_to="none",
    save_only_model=True,      # 只保存模型, 不保存 optimizer/scheduler
    save_safetensors=True,    # safetensors 更安全更快
)

def model_init():
    return AutoModelForSequenceClassification.from_pretrained(
        teacher_ckpt, num_labels=num_labels, id2label=id2label, label2id=label2id
    )

teacher_trainer = Trainer(
    model_init=model_init,
    args=teacher_training_args,
    train_dataset=clinc_enc['train'],
    eval_dataset=clinc_enc['validation'],
    compute_metrics=compute_metrics,
    tokenizer=teacher_tokenizer,
)
teacher_trainer.train()

save_path = "/kaggle/working/ModernBERT-large-finetuned-clinc"
teacher_trainer.save_model(save_path)
teacher_tokenizer.save_pretrained(save_path) # 确保分词器也保存了

```

```
/tmp/ipython-input-1078479289.py:25: FutureWarning: `tokenizer` is deprecated and will be removed in version 5.0.0 for `Trainer.__init__`. Use
teacher_trainer = Trainer(
config.json:    1.19k/? [00:00<00:00, 78.5kB/s]

model.safetensors: 100%                                1.58G/1.58G [00:41<00:00, 35.9MB/s]
Some weights of ModernBertForSequenceClassification were not initialized from the model checkpoint at answerdotai/ModernBERT-large and are not
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
The tokenizer has new PAD/BOS/EOS tokens that differ from the model config and generation config. The model config and generation config were
Some weights of ModernBertForSequenceClassification were not initialized from the model checkpoint at answerdotai/ModernBERT-large and are not
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
W1007 18:16:04.418000 1380 torch/_inductor/utils.py:1436] [1/0_1] Not enough SMs to use max_autotune_gemm mode
[4770/4770 28:31, Epoch 5/5]

Epoch Training Loss Validation Loss Accuracy
-----
```

Epoch	Training Loss	Validation Loss	Accuracy
1	0.168600	0.218951	0.952258
2	0.038900	0.185589	0.967742
3	0.018800	0.167427	0.969032
4	0.003300	0.165662	0.970323
5	0.000100	0.164647	0.969677

```
(' /kaggle/working/ModernBERT-large-finetuned-clinc/tokenizer_config.json',
 '/kaggle/working/ModernBERT-large-finetuned-clinc/special_tokens_map.json',
 '/kaggle/working/ModernBERT-large-finetuned-clinc/tokenizer.json')
```

```
# 用 trainer 评估编码后的测试集
def tokenize_with_labels(batch):
    tokens = teacher_tokenizer(batch["text"], truncation=True, padding="max_length", max_length=128)
    tokens["labels"] = batch["intent"]
    return tokens

clinc_enc_test = clinc["test"].map(tokenize_with_labels, batched=True)

trainer_metrics = teacher_trainer.evaluate(eval_dataset=clinc_enc_test)
print(trainer_metrics)
```

```
Map: 100%                                5500/5500 [00:02<00:00, 2809.07 examples/s]
[344/344 27:03]
{'eval_loss': 0.5416863560676575, 'eval_accuracy': 0.9003636363636364, 'eval_runtime': 139.9903, 'eval_samples_per_second': 39.288, 'eval_steps': 1}
```

```
print(clinc["test"][0]["text"], clinc["test"][0]["intent"])
print(id2label[clinc["test"][0]["intent"]])
```

```
how would you say fly in italian 61
translate
```

```
from huggingface_hub import login
from google.colab import userdata

# 从 Colab Secrets 中读取 HF_TOKEN
HF_TOKEN = userdata.get('colabb')

# 使用令牌登录
if HF_TOKEN:
    login(token=HF_TOKEN)
    print("Successfully logged in to Hugging Face using Colab Secret!")
else:
    print("HF_TOKEN is not set in Colab Secrets.")

Successfully logged in to Hugging Face using Colab Secret!
```

```
teacher_trainer.model.push_to_hub("pine-cone/ModernBERT-large-finetuned-clinc")
teacher_tokenizer.push_to_hub("pine-cone/ModernBERT-large-finetuned-clinc")
```

```
NameError                                 Traceback (most recent call last)
/tmp/ipython-input-4054812914.py in <cell line: 0>()
      1 teacher_trainer.model.push_to_hub("pine-cone/ModernBERT-large-finetuned")
      2
      3 teacher_tokenizer.push_to_hub("pine-cone/ModernBERT-large-finetuned")
```

```
NameError: name 'teacher_trainer' is not defined
```

```

pipe_teacher = pipeline(
    "text-classification",
    model="pine-cone/ModernBERT-large-finetuned-clinc",
    tokenizer="pine-cone/ModernBERT-large-finetuned-clinc",
)
query = "Hey, I'd like to rent a vehicle from Nov 1st to Nov 15th in Paris and I need a 15 passenger van"
print(pipe_teacher(query))

```

```

config.json:      9.17k/? [00:00<00:00, 1.01MB/s]
model.safetensors: 100%                                1.58G/1.58G [01:23<00:00, 50.8MB/s]
tokenizer_config.json:     20.8k/? [00:00<00:00, 2.14MB/s]
tokenizer.json:      3.58M/? [00:00<00:00, 11.3MB/s]
special_tokens_map.json: 100%                                694/694 [00:00<00:00, 82.2kB/s]

Device set to use cuda:0
[{'label': 'car_rental', 'score': 0.9993919134140015}]

```

```

# %% Distillation trainer
class DistillationTrainingArguments(TrainingArguments):
    def __init__(self, *args, alpha=0.5, temperature=2.0, **kwargs):
        super().__init__(*args, **kwargs)
        self.alpha, self.temperature = alpha, temperature
class DistillationTrainer(Trainer):
    def __init__(self, *args, teacher_model=None, **kwargs):
        super().__init__(*args, **kwargs)
        self.teacher_model = teacher_model
        if self.teacher_model:
            self.teacher_model.to(self.args.device); self.teacher_model.eval()
    def compute_loss(self, model, inputs, return_outputs=False, **kwargs):
        outputs_stu = model(**inputs); loss_ce, logits_stu = outputs_stu.loss, outputs_stu.logits
        with torch.no_grad(): logits_tea = self.teacher_model(**inputs).logits
        loss_fct = nn.KLDivLoss(reduction="batchmean")
        loss_kd = self.args.temperature**2 * loss_fct(
            F.log_softmax(logits_stu/self.args.temperature, dim=-1),
            F.softmax(logits_tea/self.args.temperature, dim=-1)
        )
        loss = self.args.alpha*loss_ce + (1.-self.args.alpha)*loss_kd
        return (loss, outputs_stu) if return_outputs else loss

```

```

# 设置蒸馏训练参数（固定超参数）
teacher_model = AutoModelForSequenceClassification.from_pretrained("pine-cone/ModernBERT-large-finetuned-clinc")
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
# !!! 关键步骤：移动到 GPU !!!
teacher_model.to(device)
teacher_model.eval() # 确保它处于评估模式

fixed_args = DistillationTrainingArguments(
    output_dir="/kaggle/working/bert-base-distilled-fixed",
    eval_strategy="epoch",
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
    learning_rate=2e-5,
    num_train_epochs=4,
    weight_decay=0.01,
    save_strategy="epoch",
    load_best_model_at_end=True,
    alpha=0.5,                                     # 固定蒸馏权重
    temperature=7.0,                               # 固定温度
    report_to="none",
    save_only_model=True,
    logging_strategy="steps",
    logging_steps=10,
    gradient_accumulation_steps=2
)

def model_init():
    model = AutoModelForSequenceClassification.from_pretrained(
        student_ckpt,
        num_labels=teacher_model.config.num_labels
    )
    # 对齐教师模型标签
    model.config.id2label = teacher_model.config.id2label
    model.config.label2id = teacher_model.config.label2id
    return model

# 创建蒸馏 Trainer
fixed_trainer = DistillationTrainer(
    model_init=model_init,

```

```

        teacher_model=teacher_model,
        args=fixed_args,
        train_dataset=clinc_enc_student["train"],
        eval_dataset=clinc_enc_student["validation"],
        compute_metrics=compute_metrics,
        tokenizer=student_tokenizer,
    )

# 开始训练
fixed_trainer.train()

config.json: 9.17k? [00:00<00:00, 779kB/s]

model.safetensors: 100% 1.58G/1.58G [01:00<00:00, 71.2MB/s]
/tmp/ipython-input-1428946420.py:8: FutureWarning: `tokenizer` is deprecated and will be removed in version 5.0.0 for `DistillationTrainer._super().__init__(*_args, **kwargs)
model.safetensors: 100% 440M/440M [00:17<00:00, 17.5MB/s]

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized:
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized:
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
[1908/1908 12:13, Epoch 4/4]

Epoch Training Loss Validation Loss Accuracy
-----  

1 2.710600 2.612303 0.647742  

2 2.089700 2.003171 0.887419  

3 1.757900 1.719881 0.924839  

4 1.680700 1.636488 0.934839

TrainOutput(global_step=1908, training_loss=2.275734621023982, metrics={'train_runtime': 733.9495, 'train_samples_per_second': 83.112, 'train_steps_per_second': 2.6, 'total_flos': 575587515554556.0, 'train_loss': 2.275734621023982, 'epoch': 4.0})

```

```

# 用 trainer 评估编码后的测试集
clinc_enc_test = clinc["test"].map(
    lambda x: student_tokenizer(x["text"], truncation=True, padding="max_length", max_length=128),
    batched=True
)
trainer_metrics = fixed_trainer.evaluate(eval_dataset=clinc_enc_test)
print(trainer_metrics)

```

```

# 保存模型
output_path = "/kaggle/working/bert-base-distilled-fixed"
fixed_trainer.save_model(output_path)
fixed_trainer.tokenizer.save_pretrained(output_path)

```

```

Trainer.tokenizer is now deprecated. You should use Trainer.processing_class instead.
('/kaggle/working/bert-base-distilled-fixed/tokenizer_config.json',
 '/kaggle/working/bert-base-distilled-fixed/special_tokens_map.json',
 '/kaggle/working/bert-base-distilled-fixed/vocab.txt',
 '/kaggle/working/bert-base-distilled-fixed/added_tokens.json',
 '/kaggle/working/bert-base-distilled-fixed/tokenizer.json')

```

```

model = AutoModelForSequenceClassification.from_pretrained(output_path)
tokenizer = AutoTokenizer.from_pretrained(output_path)

model.push_to_hub("pine-cone/bert-base-distilled-fixed")
tokenizer.push_to_hub("pine-cone/bert-base-distilled-fixed")

```

```

Processing Files (1 / 1) : 100% 438MB / 438MB, 135MB/s
New Data Upload : 0.00B / 0.00B, 0.00B/s
...z29e63t/model.safetensors: 100% 438MB / 438MB

No files have been modified since last commit. Skipping to prevent empty commit.
WARNING:huggingface_hub.hf_api:No files have been modified since last commit. Skipping to prevent empty commit.
No files have been modified since last commit. Skipping to prevent empty commit.
WARNING:huggingface_hub.hf_api:No files have been modified since last commit. Skipping to prevent empty commit.
CommitInfo(commit_url='https://huggingface.co/pine-cone/bert-base-distilled-fixed/commit/ba9819dd8a2faedd06c67f66a55eb846833b75ef',
commit_message='Upload tokenizer', commit_description='', oid='ba9819dd8a2faedd06c67f66a55eb846833b75ef', pr_url=None,
repo_url=RepoUrl('https://huggingface.co/pine-cone/bert-base-distilled-fixed', endpoint='https://huggingface.co', repo_type='model',
repo_id='pine-cone/bert-base-distilled-fixed') or revision=None or num=None)

```

```

# %% Distillation with Optuna
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

teacher_model = AutoModelForSequenceClassification.from_pretrained("pine-cone/ModernBERT-large-finetuned-clinc")

# !!! 关键步骤: 移动到 GPU !!!

```

```
teacher_model.to(device)
teacher_model.eval() # 确保它处于评估模式
```

显示隐藏的输出项

```
# % Distillation with Optuna

def hp_space(trial):
    return {
        "learning_rate": 2e-5, # 固定
        "num_train_epochs": 8, # 固定
        "num_train_epochs": trial.suggest_int("num_train_epochs", 5, 10),
        "alpha": trial.suggest_float("alpha", 0.4, 0.8),
        "alpha": trial.suggest_categorical("alpha", [0.5, 0.6, 0.8]),
        "temperature": trial.suggest_categorical("temperature", [4, 6, 7]),
        "temperature": trial.suggest_float("temperature", 2.0, 12.0)
    }

distil_args = DistillationTrainingArguments(
    output_dir="/kaggle/working/bert-base-distilled-Optuna",
    eval_strategy="epoch",
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
    learning_rate=2e-5,
    num_train_epochs=6, # 这个会被 Optuna 搜索覆盖
    weight_decay=0.01,
    save_strategy="epoch",
    load_best_model_at_end=True,
    metric_for_best_model="accuracy", # 明确最佳模型的评价指标
    greater_is_better=True, # accuracy 越大越好
    alpha=0.5, # 初始值，会被 Optuna 搜索覆盖
    temperature=2.0, # 初始值，会被 Optuna 搜索覆盖
    report_to="none",
    save_only_model=True,
    logging_strategy="steps",
    logging_steps=10,
    gradient_accumulation_steps=2
)

def model_init():
    model = AutoModelForSequenceClassification.from_pretrained(
        student_ckpt,
        num_labels=teacher_model.config.num_labels
    )
    # 对齐教师模型标签
    model.config.id2label = teacher_model.config.id2label
    model.config.label2id = teacher_model.config.label2id
    return model

search_trainer = DistillationTrainer(
    model_init=model_init,
    teacher_model=teacher_model,
    args=distil_args,
    train_dataset=clinc_enc_student["train"],
    eval_dataset=clinc_enc_student["validation"],
    compute_metrics=compute_metrics,
    tokenizer=student_tokenizer,
)
best_run = search_trainer.hyperparameter_search(n_trials=6, direction="maximize", hp_space=hp_space)
```

```
/tmp/ipython-input-1428946420.py:8: FutureWarning: `tokenizer` is deprecated and will be removed in version 5.0.0 for `DistillationTrainer`.
super().__init__(*args, **kwargs)

model.safetensors: 100%
440M/440M [00:08<00:00, 84.4MB/s]

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
[I 2025-10-05 12:13:14, 458] A new study created in memory with name: no-name-9e28c3db-b767-43dc-be72-209ae6ef33a8
Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
W1005 12:13:23.144000 1564 torch/_inductor/utils.py:1436] [1/0_1] Not enough SMs to use max_autotune_gemm mode
[3816/3816 22:22, Epoch 8/8]
```

Epoch	Training Loss	Validation Loss	Accuracy
1	2.675100	2.523523	0.801290
2	1.688200	1.620851	0.926452
3	1.211600	1.195066	0.945806
4	0.989000	1.011120	0.950968
5	0.865800	0.928689	0.951935
6	0.798600	0.891946	0.952903
7	0.784700	0.874794	0.954194
8	0.775000	0.868798	0.954194

```
[I 2025-10-05 12:35:48, 207] Trial 0 finished with value: 0.9541935483870968 and parameters: {'alpha': 0.630571588388972, 'temperature': 10}
Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
[3816/3816 22:19, Epoch 8/8]
```

Epoch	Training Loss	Validation Loss	Accuracy
1	2.704400	2.551690	0.810323
2	1.696400	1.628750	0.925806
3	1.208000	1.192352	0.947097
4	0.986700	1.004799	0.952581
5	0.854700	0.919138	0.954516
6	0.786600	0.880504	0.952581
7	0.773600	0.862128	0.953548
8	0.763700	0.856140	0.953548

```
[I 2025-10-05 12:58:09, 627] Trial 1 finished with value: 0.9535483870967741 and parameters: {'alpha': 0.6412513943521345, 'temperature': 10}
Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
[3816/3816 22:42, Epoch 8/8]
```

Epoch	Training Loss	Validation Loss	Accuracy
1	2.788600	2.621471	0.790323

```
#best_params = best_run.hyperparameters
def model_init():
    model = AutoModelForSequenceClassification.from_pretrained(
        student_ckpt,
        num_labels=teacher_model.config.num_labels
    )
    # 对齐教师模型标签
    model.config.id2label = teacher_model.config.id2label
    model.config.label2id = teacher_model.config.label2id
    return model

# Step 1: 创建新的训练参数对象, 填入最优 alpha / T
best_args = DistillationTrainingArguments(
    output_dir="/kaggle/working/bert-base-distilled-final",
    eval_strategy="epoch",
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
    learning_rate=2e-5,                                     # 这个是固定值 (如果没调)
    num_train_epochs=8,                                     # 或者你也可以用 best_params 里的
    weight_decay=0.01,
    warmup_ratio=0.06,
    save_strategy="epoch",
    load_best_model_at_end=True,
    metric_for_best_model="eval_macro_f1",
    greater_is_better=True,
    report_to="none",
    logging_strategy="steps",
    logging_steps=10,
    gradient_accumulation_steps=2,
```

```

# 🔥 用最优超参覆盖
alpha=0.724672262749978, #best_params["alpha"],
temperature=7#best_params["temperature"]
)

# ✅ Step 2: 初始化模型
best_model = model_init()

# ✅ Step 3: 新 Trainer(最终训练)
final_trainer = DistillationTrainer(
    model=best_model,
    teacher_model=teacher_model,
    args=best_args,
    train_dataset=clinc_enc_student["train"],
    eval_dataset=clinc_enc_student["validation"],
    compute_metrics=compute_metrics,
    tokenizer=student_tokenizer
)

# ✅ Step 4: 重新训练(用最优超参)
final_trainer.train()

```

```

model.safetensors: 100%                                         440M/440M [00:05<00:00, 101MB/s]
Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized:
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
/tmp/ipython-input-1428946420.py:8: FutureWarning: `tokenizer` is deprecated and will be removed in version 5.0.0 for `DistillationTrainer._super().__init__(*_args, **kwargs)
W1008 12:35:31.063000 1475 torch/_inductor/utils.py:1436] [1/0_1] Not enough SMs to use max_autotune_gemm mode
[3816/3816 26:35, Epoch 8/8]

Epoch Training Loss Validation Loss Accuracy Macro F1
-----  

1      2.987700      2.848179      0.760968      0.749750  

2      1.685000      1.608866      0.914194      0.917152  

3      1.122100      1.092647      0.943871      0.948558  

4      0.862700      0.886773      0.950645      0.955031  

5      0.718800      0.798898      0.953548      0.958233  

6      0.641000      0.758899      0.954194      0.958709  

7      0.632800      0.737825      0.955161      0.959764  

8      0.626700      0.733441      0.953871      0.958537

```

TrainOutput(global_step=3816, training_loss=1.3709464952880732, metrics={'train_runtime': 1605.7092, 'train_samples_per_second': 75.979, 'train_steps_per_second': 2.377, 'total_flos': 1149331818924300.0, 'train_loss': 1.3709464952880732, 'epoch': 8.0})

```

# 用 trainer 评估编码后的测试集
def tokenize_with_labels(batch):
    tokens = student_tokenizer(batch["text"], truncation=True, padding="max_length", max_length=128)
    tokens["labels"] = batch["intent"]
    return tokens

clinc_enc_test = clinc["test"].map(tokenize_with_labels, batched=True)

trainer_metrics = final_trainer.evaluate(eval_dataset=clinc_enc_test)
print(trainer_metrics)

```

```

Map: 100%                                         5500/5500 [00:01<00:00, 4524.54 examples/s]
[344/344 03:15]
{'eval_loss': 0.8634978532791138, 'eval_accuracy': 0.9056363636363637, 'eval_macro_f1': 0.9282288173816916, 'eval_runtime': 196.1271, 'eval_s'

```

```

output_path = "/kaggle/working/bert-base-distilled-final"
final_trainer.save_model(output_path)
final_trainer.tokenizer.save_pretrained(output_path)

```

```

Trainer.tokenizer is now deprecated. You should use Trainer.processing_class instead.
(' /kaggle/working/bert-base-distilled-final/tokenizer_config.json',
 '/kaggle/working/bert-base-distilled-final/special_tokens_map.json',
 '/kaggle/working/bert-base-distilled-final/vocab.txt',
 '/kaggle/working/bert-base-distilled-final/added_tokens.json',
 '/kaggle/working/bert-base-distilled-final/tokenizer.json')

```

```

final_model = final_trainer.model
final_tokenizer = student_tokenizer

final_model.push_to_hub("pine-cone/bert-base-distilled-Optuna")
final_tokenizer.push_to_hub("pine-cone/bert-base-distilled-Optuna")
#output_path = best_run.checkpoint
#model = AutoModelForSequenceClassification.from_pretrained(output_path)
#tokenizer = AutoTokenizer.from_pretrained(output_path)

#model.push_to_hub("pine-cone/bert-base-distilled-Optuna")
#tokenizer.push_to_hub("pine-cone/bert-base-distilled-Optuna")

README.md: 5.17k? [00:00<00:00, 173kB/s]

Processing Files (1 / 1) : 100% 438MB / 438MB, 33.3MB/s

New Data Upload : 100% 438MB / 438MB, 33.3MB/s

..v8wh1ci/model.safetensors: 100% 438MB / 438MB

CommitInfo(commit_url='https://huggingface.co/pine-cone/bert-base-distilled-Optuna/commit/b3a74bc2360c2e9184b19e47e28f177f87d0a0e1', commit_message='Upload tokenizer', commit_description='', oid='b3a74bc2360c2e9184b19e47e28f177f87d0a0e1', pr_url=None, repo_url=RepoUrl('https://huggingface.co/pine-cone/bert-base-distilled-Optuna', endpoint='https://huggingface.co', repo_type='model', rev_id='pine-cone/bert-base-distilled-Optuna'), revision=None, num=None)

# %% Benchmark class
clinc = load_dataset("clinc_oos", "plus")
test_ds = clinc["test"]
intents = test_ds.features["intent"]
accuracy_score = evaluate.load("accuracy")

class PerformanceBenchmark:
    def __init__(self, pipeline, dataset, optim_type):
        self.pipeline, self.dataset, self.optim_type = pipeline, dataset, optim_type
    def compute_accuracy(self):
        preds, labels = [], []
        for example in self.dataset:
            pred = self.pipeline(example["text"])[0]["label"]
            #preds.append(intents.str2int(pred)); labels.append(example["intent"])
            pred_label = pred.replace("LABEL_", "")
            pred_idx = int(pred_label) if pred_label.isdigit() else intents.str2int(pred_label)
            preds.append(pred_idx)
            labels.append(example["intent"])

        return accuracy_score.compute(predictions=preds, references=labels)
    def compute_size(self):
        state_dict = self.pipeline.model.state_dict()
        tmp_path = Path("model.pt"); torch.save(state_dict, tmp_path)
        size_mb = Path(tmp_path).stat().st_size/(1024*1024); tmp_path.unlink()
        return {"size_mb": size_mb}
    def time_pipeline(self, query="What is the pin number for my account?"):
        for _ in range(10): _ = self.pipeline(query)
        latencies = []
        for _ in range(100):
            start = perf_counter(); _ = self.pipeline(query); latencies.append(perf_counter()-start)
        return {"time_avg_ms": 1000*np.mean(latencies), "time_std_ms": 1000*np.std(latencies)}
    def run_benchmark(self):
        m = {}; m[self.optim_type] = {}
        m[self.optim_type].update(self.compute_size())
        m[self.optim_type].update(self.time_pipeline())
        m[self.optim_type].update(self.compute_accuracy())
        return m

perf_metrics = {}

model_path = "pine-cone/ModernBERT-large-finetuned-clinc"
model = AutoModelForSequenceClassification.from_pretrained(model_path)
model.eval()
pipe_teacher_cpu = pipeline(
    "text-classification",
    model=model,
    tokenizer=model_path,
    device=-1 # CPU
)

pb = PerformanceBenchmark(pipe_teacher_cpu, clinc["test"], "ModernBERT-large")
perf_metrics.update(pb.run_benchmark())

```

```
tokenizer_config.json: 20.8k/? [00:00<00:00, 1.98MB/s]
tokenizer.json: 3.58M/? [00:00<00:00, 103MB/s]
special_tokens_map.json: 100%
Device set to use cpu
```

```
# 推理管道与性能评估

model_path = "pine-cone/bert-base-distilled-fixed"
model = AutoModelForSequenceClassification.from_pretrained(model_path)
model.eval()
pipe_fixed_cpu = pipeline(
    "text-classification",
    model=model,
    tokenizer=model_path,
    device=-1
)

pb = PerformanceBenchmark(pipe_fixed_cpu, clinc["test"], "Distillation-Fixed")
perf_metrics.update(pb.run_benchmark())
```

```
config.json: 8.27k/? [00:00<00:00, 544kB/s]
model.safetensors: 100% 438M/438M [00:21<00:00, 24.8MB/s]
tokenizer_config.json: 1.41k/? [00:00<00:00, 123kB/s]
vocab.txt: 232k/? [00:00<00:00, 9.29MB/s]
tokenizer.json: 712k/? [00:00<00:00, 12.7MB/s]
special_tokens_map.json: 100% 695/695 [00:00<00:00, 22.3kB/s]
Device set to use cpu
```

```
model_path = "pine-cone/bert-base-distilled-final"
model = AutoModelForSequenceClassification.from_pretrained(model_path)
model.eval()
pipe_distilled_cpu = pipeline(
    "text-classification",
    model=model,
    tokenizer=model_path,
    device=-1
)

pb = PerformanceBenchmark(pipe_distilled_cpu, clinc["test"], "Distillation-Optuna")
perf_metrics.update(pb.run_benchmark())
```

```
config.json: 8.57k/? [00:00<00:00, 791kB/s]
model.safetensors: 100% 438M/438M [00:16<00:00, 29.4MB/s]
tokenizer_config.json: 1.22k/? [00:00<00:00, 52.5kB/s]
vocab.txt: 232k/? [00:00<00:00, 6.86MB/s]
tokenizer.json: 711k/? [00:00<00:00, 9.66MB/s]
special_tokens_map.json: 100% 125/125 [00:00<00:00, 11.9kB/s]
Device set to use cpu
```

```
# %% Quantization
model_name = "pine-cone/bert-base-distilled-Optuna"
output_path = "/kaggle/working/bert-base-distilled-final"
model_distilled = AutoModelForSequenceClassification.from_pretrained(output_path).to("cpu")
model_distilled.eval()
model_quantized = quantize_dynamic(model_distilled, {torch.nn.Linear}, dtype=torch.qint8)

pipe_quantized_cpu = pipeline(
    "text-classification",
    model=model_quantized,
    tokenizer=output_path,
    device=-1
)
model.eval()

pb = PerformanceBenchmark(pipe_quantized_cpu, clinc["test"], "Distillation+Quantization")
perf_metrics.update(pb.run_benchmark())
```

```
/tmp/ipython-input-2551072981.py:6: DeprecationWarning: torch.ao.quantization is deprecated and will be removed in 2.10.
For migrations of users:
```

1. Eager mode quantization (torch.ao.quantization.quantize, torch.ao.quantization.quantize_dynamic), please migrate to use torchao eager mode
 2. FX graph mode quantization (torch.ao.quantization.quantize_fx.prepare_fx, torch.ao.quantization.quantize_fx.convert_fx, please migrate to use torchao eager mode)
 3. pt2e quantization has been migrated to torchao (<https://github.com/pytorch/ao/tree/main/torchao/quantization/pt2e>)
 see <https://github.com/pytorch/ao/issues/2259> for more details
 model_quantized = quantize_dynamic(model_distilled, {torch.nn.Linear}, dtype=torch.qint8)
 Device set to use cpu

```

# %% Plot results
final_df = pd.DataFrame.from_dict(perf_metrics, orient='index')
print(final_df)

      size_mb  time_avg_ms  time_std_ms  accuracy
ModernBERT-large      1510.624507     521.414343     137.759554  0.915091
Distillation-Fixed      418.147475     116.434022      9.250822  0.902182
Distillation-Optuna      418.147475     117.118999      8.689499  0.907818
Distillation+Quantization   173.182982      30.170859      2.525748  0.904182
  
```

```

%matplotlib inline
print(perf_metrics)

def plot_metrics(perf_metrics, current_optim_type=None):
    df = pd.DataFrame.from_dict(perf_metrics, orient='index')

    for idx in df.index:
        df_opt = df.loc[idx]
        if idx == current_optim_type:
            plt.scatter(df_opt["time_avg_ms"], df_opt["accuracy"] * 100,
                        alpha=0.5, s=df_opt["size_mb"], label=idx,
                        marker='$\u25CC$')
        else:
            plt.scatter(df_opt["time_avg_ms"], df_opt["accuracy"] * 100,
                        s=df_opt["size_mb"], label=idx, alpha=0.5)

    plt.legend(bbox_to_anchor=(1, 1))
    plt.ylabel("Accuracy (%)")
    plt.xlabel("Average Latency (ms)")
    plt.title("Accuracy vs Latency (bubble size = model size)")
    plt.ylim(85, 95) # 或自动根据结果范围设定
    plt.grid(True)
    plt.show()

plot_metrics(perf_metrics, current_optim_type="DistilBERT")
  
```

```
{'ModernBERT-large': {'size_mb': 1510.6245069503784, 'time_avg_ms': np.float64(521.4143433500021), 'time_std_ms': np.float64(137.759553569851}
```